

PART A — Email Tagging System

1. Problem Statement:

- The goal of Part A is to design an email-tagging system that automatically assigns issue tags to incoming customer emails. The system must
 - Work across multiple customers
 - Avoid cross-customer tag leakage
 - Perform well despite limited training examples
 - Use a combination of ML, rules, and guardrails
 - Provide a reliable and explainable tagging workflow
-

2. Dataset Overview:

- Two datasets were provided:
 - small_dataset.csv (very small, only a few samples per tag)
 - large_dataset.csv (more tags but still very sparse)
- Dataset characteristics:
 - ~60 unique tags
 - Most tags appear only once
 - Per-customer data is extremely small (1–3 samples)
 - Text content is extremely short and sparse

This makes training a pure machine-learning model very difficult, especially for multi-class classification.

3. Baseline Model (TF-IDF + Logistic Regression):

- A baseline multi-class classifier was trained using:
 - TF-IDF features
 - Logistic Regression (max_iter=2000)
 - Train/Test split or full-dataset training depending on size
- Due to extreme sparsity of labels and low sample count:

- Model predictions collapse to a few classes
- F1, precision, and recall across most classes are 0
- Confusion matrix shows most predictions along one column
- Model cannot generalize

Key Reason for Poor Performance

- Most classes have only 1 sample
- No meaningful patterns can be learned
- Vocabulary overlap is extremely low

This is expected and an important part of the assignment.

4. Error Analysis (Why the Model Fails)

Observations

1. Many tags have only one example → impossible to learn boundaries
2. Emails are extremely short → limited feature extraction
3. Multi-class (60 labels) with very few samples → underfitting
4. Overlap between tag semantics is high
5. Per-customer data is even smaller

Conclusion

Traditional ML cannot solve this dataset alone.

This motivates customer isolation and rule-based guardrails.

5. Customer Isolation

To prevent cross-customer tag leakage, a mapping of allowed tags per customer was created:

```
customer_allowed_tags = {  
    "CUST_A": ["workflow_issue", "notification_bug", ...],  
    "CUST_B": ["sync_bug", "analytics_issue", ...]  
}
```

When predicting tags:

- If a customer-specific model exists → use it
- If not → use the global model but **filter predictions** only to that customer's allowed tags

Impact

- Reduces prediction space from ~60 tags to 3–5 tags per customer
 - Increases accuracy
 - Prevents invalid tag assignment
 - Automatically improves reliability
-

6. Customer-Specific Models

Per-customer models were trained **only if a customer had ≥ 5 samples**.

Benefits:

- Much smaller number of tags
- Local patterns per customer can be learned
- Influences the final prediction pipeline

Even if accuracy remains low, customer-specific models still narrow down predictions.

7. Guardrails (Pattern + Anti-Pattern Rules)

Because ML alone is insufficient, guardrails were added to correct/override predictions.

Sample Guardrail Rules

Billing Issues

if text contains: charge, invoice, billed, refund
→ "billing_error"

Workflow / Rules / Automation

if text contains: rule, workflow, sla, automation
→ "workflow_issue"

Notifications

if text contains: notification, alert
→ "notification_bug"

Mobile Issues

if text contains: mobile
→ "mobile_bug"

Tagging Issues

if text contains: tag, tagging
→ "tagging_issue"
Guardrails fix many incorrect model predictions.

8. Final Prediction Pipeline

The final pipeline combines:

1 Customer-specific model

If available → used directly.

2 Global model with tag filtering

Only tags allowed for that customer are considered.

3 Guardrail corrections

Overrides ML predictions using rule patterns.

4 Return structured output

```
{  
    predicted: <model prediction>,  
    corrected: <guardrail prediction>,  
    confidence: <score>,  
    source: "customer_model" | "global_filtered"  
}
```

9. Example Output

Input:

"We are unable to configure auto assignment rules."

ML predicted:

auth_issue

Guardrails corrected it to:

workflow_issue

Final output:

```
{  
  'predicted': 'auth_issue',  
  'corrected': 'workflow_issue',  
  'confidence': 1.0,  
  'source': 'customer_model'  
}
```

This demonstrates how ML + rules improve reliability.

10. Confusion Matrix Interpretation

The confusion matrix shows:

- Most cells = zero
 - A few repeated predictions
 - Model collapses to a few tags
 - Expected due to sparse training data
 - Motivates rule-based corrections and customer isolation
-

11. System Architecture Diagram

(Describe in Word, or I can generate diagram text)

Flow:

Email → Preprocessing → Customer Model →
→ (if unavailable → Global Model + Tag Filtering) →
→ Guardrails → Final Tag Output

12. Production Improvements

1. **Use embeddings (BERT/SentenceTransformers) instead of TF-IDF**
Better semantic understanding.
 2. **Add contextual RAG-based retrieval**
Fetch similar past tickets and use similarity to improve prediction.
 3. **Human-in-the-loop feedback loop**
Collect agent corrections → retrain periodically.
 4. **Hybrid ML + Rule Engine**
Rules for high-precision domains, ML for general cases.
-

13. Final Deliverables (for Word/PDF)

Include:

- Baseline model explanation
- Dataset summary
- Error analysis
- Customer isolation logic
- Guardrail rules
- Final prediction architecture
- Confusion matrix screenshot
- Sample output
- Production improvements
- Conclusion

